

Residual Networks

Welcome to the first assignment of this week! You'll be building a very deep convolutional network, using Residual Networks (ResNets). In theory, very deep networks can represent very complex functions; but in practice, they are hard to train. Residual Networks, introduced by [He et al. \(https://arxiv.org/pdf/1512.03385.pdf\)](https://arxiv.org/pdf/1512.03385.pdf), allow you to train much deeper networks than were previously feasible.

By the end of this assignment, you'll be able to:

- Implement the basic building blocks of ResNets in a deep neural network using Keras
- Put together these building blocks to implement and train a state-of-the-art neural network for image classification
- Implement a skip connection in your network

For this assignment, you'll use Keras.

Before jumping into the problem, run the cell below to load the required packages.

Table of Content

- [1 - Packages](#)
- [2 - The Problem of Very Deep Neural Networks](#)
- [3 - Building a Residual Network](#)
 - [3.1 - The Identity Block](#)
 - [Exercise 1 - identity_block](#)
 - [3.2 - The Convolutional Block](#)
 - [Exercise 2 - convolutional_block](#)
- [4 - Building Your First ResNet Model \(50 layers\)](#)
 - [Exercise 3 - ResNet50](#)
- [5 - Test on Your Own Image \(Optional/Ungraded\)](#)
- [6 - Bibliography](#)

1 - Packages

```
In [1]: import tensorflow as tf
import numpy as np
import scipy.misc
from tensorflow.keras.applications.resnet_v2 import ResNet50V2
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet_v2 import preprocess_input, decode_predictions
from tensorflow.keras import layers
from tensorflow.keras.layers import Input, Add, Dense, Activation, ZeroPadding2D, BatchNormalization, Flatten, Conv2D, AveragePooling2D, MaxPooling2D, GlobalMaxPooling2D
from tensorflow.keras.models import Model, load_model
from resnets_utils import *
from tensorflow.keras.initializers import random_uniform, glorot_uniform, constant, identity
from tensorflow.python.framework.ops import EagerTensor
from matplotlib.pyplot import imshow

from test_utils import summary, comparator
import public_tests

%matplotlib inline
```

2 - The Problem of Very Deep Neural Networks

Last week, you built your first convolutional neural networks: first manually with numpy, then using Tensorflow and Keras.

In recent years, neural networks have become much deeper, with state-of-the-art networks evolving from having just a few layers (e.g., AlexNet) to over a hundred layers.

- The main benefit of a very deep network is that it can represent very complex functions. It can also learn features at many different levels of abstraction, from edges (at the shallower layers, closer to the input) to very complex features (at the deeper layers, closer to the output).
- However, using a deeper network doesn't always help. A huge barrier to training them is vanishing gradients: very deep networks often have a gradient signal that goes to zero quickly, thus making gradient descent prohibitively slow.
- More specifically, during gradient descent, as you backpropagate from the final layer back to the first layer, you are multiplying by the weight matrix on each step, and thus the gradient can decrease exponentially quickly to zero (or, in rare cases, grow exponentially quickly and "explode," from gaining very large values).
- During training, you might therefore see the magnitude (or norm) of the gradient for the shallower layers decrease to zero very rapidly as training proceeds, as shown below:




Figure 1 : Vanishing gradient
The speed of learning decreases very rapidly for the shallower layers as the network trains

Not to worry! You are now going to solve this problem by building a Residual Network!

3 - Building a Residual Network

In ResNets, a "shortcut" or a "skip connection" allows the model to skip layers:




Figure 2: A ResNet block showing a skip-connection

The image on the left shows the "main path" through the network. The image on the right adds a shortcut to the main path. By stacking these ResNet blocks on top of each other, you can form a very deep network. The lecture mentioned that having ResNet blocks with the shortcut also makes it very easy for one of the blocks to learn an identity function. This means that you can stack on additional ResNet blocks with little risk of harming training set performance.

On that note, there is also some evidence that the ease of learning an identity function accounts for ResNets' remarkable performance even more than skip connections help with vanishing gradients.

Two main types of blocks are used in a ResNet, depending mainly on whether the input/output dimensions are the same or different. You are going to implement both of them: the "identity block" and the "convolutional block."

3.1 - The Identity Block

The identity block is the standard block used in ResNets, and corresponds to the case where the input activation (say $a^{[l]}$) has the same dimension as the output activation (say $a^{[l+2]}$). To flesh out the different steps of what happens in a ResNet's identity block, here is an alternative diagram showing the individual steps:




Figure 3: Identity block. Skip connection "skips over" 2 layers.

The upper path is the "shortcut path." The lower path is the "main path." In this diagram, notice the CONV2D and ReLU steps in each layer. To speed up training, a BatchNorm step has been added. Don't worry about this being complicated to implement—you'll see that BatchNorm is just one line of code in Keras!

In this exercise, you'll actually implement a slightly more powerful version of this identity block, in which the skip connection "skips over" 3 hidden layers rather than 2 layers. It looks like this:




Figure 4: Identity block. Skip connection "skips over" 3 layers.

These are the individual steps:

First component of main path:

- The first CONV2D has F_1 filters of shape $(1,1)$ and a stride of $(1,1)$. Its padding is "valid". Use 0 as the seed for the random uniform initialization: `kernel_initializer = initializer(seed=0)`.
- The first BatchNorm is normalizing the 'channels' axis.
- Then apply the ReLU activation function. This has no hyperparameters.

Second component of main path:

- The second CONV2D has F_2 filters of shape (f, f) and a stride of $(1,1)$. Its padding is "same". Use 0 as the seed for the random uniform initialization: `kernel_initializer = initializer(seed=0)`.
- The second BatchNorm is normalizing the 'channels' axis.
- Then apply the ReLU activation function. This has no hyperparameters.

Third component of main path:

- The third CONV2D has F_3 filters of shape $(1,1)$ and a stride of $(1,1)$. Its padding is "valid". Use 0 as the seed for the random uniform initialization: `kernel_initializer = initializer(seed=0)`.
- The third BatchNorm is normalizing the 'channels' axis.
- Note that there is **no** ReLU activation function in this component.

Final step:

- The `x_shortcut` and the output from the 3rd layer `x` are added together.
- **Hint:** The syntax will look something like `Add()([var1, var2])`
- Then apply the ReLU activation function. This has no hyperparameters.

Exercise 1 - identity_block

Implement the ResNet identity block. The first component of the main path has been implemented for you already! First, you should read these docs carefully to make sure you understand what's happening. Then, implement the rest.

- To implement the Conv2D step: [Conv2D \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D)
- To implement BatchNorm: [BatchNormalization \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/BatchNormalization\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/BatchNormalization) `BatchNormalization(axis = 3)(X, training = training)`. If training is set to False, its weights are not updated with the new examples. I.e. when the model is used in prediction mode.
- For the activation, use: `Activation('relu')(X)`
- To add the value passed forward by the shortcut: [Add \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/Add\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Add)

We have added the `initializer` argument to our functions. This parameter receives an initializer function like the ones included in the package [tensorflow.keras.initializers](https://www.tensorflow.org/api_docs/python/tf/keras/initializers) (https://www.tensorflow.org/api_docs/python/tf/keras/initializers) or any other custom initializer. By default it will be set to `random_uniform` (https://www.tensorflow.org/api_docs/python/tf/keras/initializers/RandomUniform)

Remember that these functions accept a `seed` argument that can be any value you want, but that in this notebook must set to 0 for **grading purposes**.

Here is where you're actually using the power of the Functional API to create a shortcut path:

```
In [6]: # GRADED FUNCTION: identity_block

def identity_block(X, f, filters, training=True, initializer=random_uniform):
    """
    Implementation of the identity block as defined in Figure 4

    Arguments:
    X -- input tensor of shape (m, n_H_prev, n_W_prev, n_C_prev)
    f -- integer, specifying the shape of the middle CONV's window for the main path
    filters -- python list of integers, defining the number of filters in the CONV layers of the main path
    training -- True: Behave in training mode
                  False: Behave in inference mode
    initializer -- to set up the initial weights of a layer. Equals to random uniform initializer

    Returns:
    X -- output of the identity block, tensor of shape (n_H, n_W, n_C)
    """

    # Retrieve Filters
    F1, F2, F3 = filters

    # Save the input value. You'll need this later to add back to the main path.
    X_shortcut = X
    cache = []

    # First component of main path
    X = Conv2D(filters = F1, kernel_size = 1, strides = (1,1), padding = 'valid', kernel_initializer = initializer(seed=0))(X)
    X = BatchNormalization(axis = 3)(X, training = training) # Default axis
    X = Activation('relu')(X)

    # YOUR CODE STARTS HERE

    ## Second component of main path (~3 lines)
    X = Conv2D(filters = F2, kernel_size = f, strides = (1,1), padding = 'same', kernel_initializer = initializer(seed=0))(X)
    X = BatchNormalization(axis = 3)(X, training = training) # Default axis
    X = Activation('relu')(X)

    ## Third component of main path (~2 lines)
    X = Conv2D(filters = F3, kernel_size = 1, strides = (1,1), padding = 'valid', kernel_initializer = initializer(seed=0))(X)
    X = BatchNormalization(axis = 3)(X, training = training) # Default axis

    ## Final step: Add shortcut value to main path, and pass it through a RELU activation (~2 lines)
    X = Add()([X, X_shortcut])
    X = Activation('relu')(X)

    # YOUR CODE ENDS HERE
    return X
```

```
In [7]: np.random.seed(1)
X1 = np.ones((1, 4, 4, 3)) * -1
X2 = np.ones((1, 4, 4, 3)) * 1
X3 = np.ones((1, 4, 4, 3)) * 3

X = np.concatenate((X1, X2, X3), axis = 0).astype(np.float32)

A3 = identity_block(X, f=2, filters=[4, 4, 3],
                    initializer=lambda seed=0:constant(value=1),
                    training=False)
A3np = A3.numpy()
print(np.around(A3.numpy()[:(0,-1),:,:].mean(axis = 3), 5))
resume = A3np[:,(0,-1),:,:].mean(axis = 3)
print(resume[1, 1, 0], resume[1, 1, 0])

public_tests.identity_block_test(identity_block)

[[[ 0.      0.      0.      0.      ]
  [ 0.      0.      0.      0.      ]

  [[192.71236 192.71236 192.71236  96.85618]
   [ 96.85618  96.85618  96.85618  48.9281 ]]]

  [[578.13715 578.13715 578.13715 290.56854]
   [290.56854 290.56854 290.56854 146.78427]]]
 96.85618 96.85618
All tests passed!
```

Expected value

```
[[[ 0.      0.      0.      0.      ]
  [ 0.      0.      0.      0.      ]

  [[192.71236 192.71236 192.71236  96.85618]
   [ 96.85618  96.85618  96.85618  48.9281 ]]]

  [[578.13715 578.13715 578.13715 290.56854]
   [290.56854 290.56854 290.56854 146.78427]]]
96.85618 96.85618
```

3.2 - The Convolutional Block

The ResNet "convolutional block" is the second block type. You can use this type of block when the input and output dimensions don't match up. The difference with the identity block is that there is a CONV2D layer in the shortcut path:




Figure 4: Convolutional block

- The CONV2D layer in the shortcut path is used to resize the input x to a different dimension, so that the dimensions match up in the final addition needed to add the shortcut value back to the main path. (This plays a similar role as the matrix W_s discussed in lecture.)
- For example, to reduce the activation dimensions's height and width by a factor of 2, you can use a 1x1 convolution with a stride of 2.
- The CONV2D layer on the shortcut path does not use any non-linear activation function. Its main role is to just apply a (learned) linear function that reduces the dimension of the input, so that the dimensions match up for the later addition step.
- As for the previous exercise, the additional `initializer` argument is required for grading purposes, and it has been set by default to `glorot_uniform` (https://www.tensorflow.org/api_docs/python/tf/keras/initializers/GlorotUniform)

The details of the convolutional block are as follows.

First component of main path:

- The first CONV2D has F_1 filters of shape (1,1) and a stride of (s,s). Its padding is "valid". Use 0 as the `glorot_uniform` seed `kernel_initializer = initializer(seed=0)`.
- The first BatchNorm is normalizing the 'channels' axis.
- Then apply the ReLU activation function. This has no hyperparameters.

Second component of main path:

- The second CONV2D has F_2 filters of shape (f,f) and a stride of (1,1). Its padding is "same". Use 0 as the `glorot_uniform` seed `kernel_initializer = initializer(seed=0)`.
- The second BatchNorm is normalizing the 'channels' axis.
- Then apply the ReLU activation function. This has no hyperparameters.

Third component of main path:

- The third CONV2D has F_3 filters of shape (1,1) and a stride of (1,1). Its padding is "valid". Use 0 as the `glorot_uniform` seed `kernel_initializer = initializer(seed=0)`.
- The third BatchNorm is normalizing the 'channels' axis. Note that there is no ReLU activation function in this component.

Shortcut path:

- The CONV2D has F_3 filters of shape (1,1) and a stride of (s,s). Its padding is "valid". Use 0 as the `glorot_uniform` seed `kernel_initializer = initializer(seed=0)`.
- The BatchNorm is normalizing the 'channels' axis.

Final step:

- The shortcut and the main path values are added together.
- Then apply the ReLU activation function. This has no hyperparameters.

Exercise 2 - convolutional_block

Implement the convolutional block. The first component of the main path is already implemented; then it's your turn to implement the rest! As before, always use 0 as the seed for the random initialization, to ensure consistency with the grader.

- `Conv2D` (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D)
- `BatchNormalization` (https://www.tensorflow.org/api_docs/python/tf/keras/layers/BatchNormalization) (axis: Integer, the axis that should be normalized (typically the features axis)) `BatchNormalization(axis = 3)(X, training = training)`. If training is set to False, its weights are not updated with the new examples. I.e when the model is used in prediction mode.
- For the activation, use: `Activation('relu')(X)`
- `Add` (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Add)

We have added the `initializer` argument to our functions. This parameter receives an initializer function like the ones included in the package `tensorflow.keras.initializers` (https://www.tensorflow.org/api_docs/python/tf/keras/initializers) or any other custom initializer. By default it will be set to `random_uniform` (https://www.tensorflow.org/api_docs/python/tf/keras/initializers/RandomUniform)

Remember that these functions accept a `seed` argument that can be any value you want, but that in this notebook must set to 0 for **grading purposes**.

```
In [8]: # GRADED FUNCTION: convolutional_block

def convolutional_block(X, f, filters, s = 2, training=True, initializer=glorot_uniform):
    """
    Implementation of the convolutional block as defined in Figure 4

    Arguments:
    X -- input tensor of shape (m, n_H_prev, n_W_prev, n_C_prev)
    f -- integer, specifying the shape of the middle CONV's window for the main path
    filters -- python list of integers, defining the number of filters in the CONV layers of the main path
    s -- Integer, specifying the stride to be used
    training -- True: Behave in training mode
                  False: Behave in inference mode
    initializer -- to set up the initial weights of a layer. Equals to Glorot uniform initializer,
                  also called Xavier uniform initializer.

    Returns:
    X -- output of the convolutional block, tensor of shape (n_H, n_W, n_C)
    """

    # Retrieve Filters
    F1, F2, F3 = filters

    # Save the input value
    X_shortcut = X

    ##### MAIN PATH #####
    # First component of main path glorot_uniform(seed=0)
    X = Conv2D(filters = F1, kernel_size = 1, strides = (s, s), padding='valid', kernel_initializer = initializer(seed=0))(X)
    X = BatchNormalization(axis = 3)(X, training=training)
    X = Activation('relu')(X)

    # YOUR CODE STARTS HERE

    ## Second component of main path (~3 lines)
    X = Conv2D(filters = F2, kernel_size = f, strides = 1, padding='same', kernel_initializer = initializer(seed=0))(X)
    X = BatchNormalization(axis = 3)(X, training=training)
    X = Activation('relu')(X)

    ## Third component of main path (~2 lines)
    X = Conv2D(filters = F3, kernel_size = 1, strides = 1, padding='valid', kernel_initializer = initializer(seed=0))(X)
    X = BatchNormalization(axis = 3)(X, training=training)

    # YOUR CODE ENDS HERE

    ##### SHORTCUT PATH ##### (~2 lines)
    # YOUR CODE STARTS HERE

    X_shortcut = Conv2D(filters = F3, kernel_size = 1, strides = s, padding='valid', kernel_initializer = initializer(seed=0))(X_shortcut)
    X_shortcut = BatchNormalization(axis = 3)(X_shortcut, training=training)

    # YOUR CODE ENDS HERE

    # Final step: Add shortcut value to main path (Use this order [X, X_shortcut]), and pass it through a RELU activation
    X = Add()([X, X_shortcut])
    X = Activation('relu')(X)

    return X
```

```
In [9]: from outputs import convolutional_block_output1
np.random.seed(1)
#X = np.random.randn(3, 4, 4, 6).astype(np.float32)
X1 = np.ones((1, 4, 4, 3)) * -1
X2 = np.ones((1, 4, 4, 3)) * 1
X3 = np.ones((1, 4, 4, 3)) * 3

X = np.concatenate((X1, X2, X3), axis = 0).astype(np.float32)

A = convolutional_block(X,
                        f = 2,
                        filters = [2, 4, 6],
                        training=False)

assert type(A) == EagerTensor, "Use only tensorflow and keras functions"
assert tuple(tf.shape(A).numpy()) == (3, 2, 2, 6), "Wrong shape. "
assert np.allclose(A.numpy(), convolutional_block_output1)
print(A[0])
print('\033[92mAll tests passed!')

tf.Tensor(
[[[0.          0.66683817 0.          0.          0.888539  0.5274254 ]
 [0.          0.65053666 0.          0.          0.8959285  0.49965227]]

 [[0.          0.6312079  0.          0.          0.86362475 0.47643146]
 [0.          0.56883204 0.          0.          0.8553412  0.417093  ]]], shape=(2, 2, 6), dtype=float32)
All tests passed!
```

Expected value

```
tf.Tensor(
[[[0.          0.66683817 0.          0.          0.888539  0.5274254 ]
 [0.          0.65053666 0.          0.          0.8959285  0.49965227]]

 [[0.          0.6312079  0.          0.          0.86362475 0.47643146]
 [0.          0.56883204 0.          0.          0.8553412  0.417093  ]]], shape=(2, 2, 6), dtype=float32)
```

4 - Building Your First ResNet Model (50 layers)

You now have the necessary blocks to build a very deep ResNet. The following figure describes in detail the architecture of this neural network. "ID BLOCK" in the diagram stands for "Identity block," and "ID BLOCK x3" means you should stack 3 identity blocks together.




Figure 5 : ResNet-50 model

The details of this ResNet-50 model are:

- Zero-padding pads the input with a pad of (3,3)
- Stage 1:
 - The 2D Convolution has 64 filters of shape (7,7) and uses a stride of (2,2).
 - BatchNorm is applied to the 'channels' axis of the input.
 - MaxPooling uses a (3,3) window and a (2,2) stride.
- Stage 2:
 - The convolutional block uses three sets of filters of size [64,64,256], "f" is 3, and "s" is 1.
 - The 2 identity blocks use three sets of filters of size [64,64,256], and "f" is 3.
- Stage 3:
 - The convolutional block uses three sets of filters of size [128,128,512], "f" is 3 and "s" is 2.
 - The 3 identity blocks use three sets of filters of size [128,128,512] and "f" is 3.
- Stage 4:
 - The convolutional block uses three sets of filters of size [256, 256, 1024], "f" is 3 and "s" is 2.
 - The 5 identity blocks use three sets of filters of size [256, 256, 1024] and "f" is 3.
- Stage 5:
 - The convolutional block uses three sets of filters of size [512, 512, 2048], "f" is 3 and "s" is 2.
 - The 2 identity blocks use three sets of filters of size [512, 512, 2048] and "f" is 3.
- The 2D Average Pooling uses a window of shape (2,2).
- The 'flatten' layer doesn't have any hyperparameters.
- The Fully Connected (Dense) layer reduces its input to the number of classes using a softmax activation.

Exercise 3 - ResNet50

Implement the ResNet with 50 layers described in the figure above. We have implemented Stages 1 and 2. Please implement the rest. (The syntax for implementing Stages 3-5 should be quite similar to that of Stage 2) Make sure you follow the naming convention in the text above.

You'll need to use this function:

- Average pooling [see reference \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/AveragePooling2D\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/AveragePooling2D)

Here are some other functions we used in the code below:

- Conv2D: [See reference \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D)
- BatchNorm: [See reference \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/BatchNormalization\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/BatchNormalization) (axis: Integer, the axis that should be normalized (typically the features axis))
- Zero padding: [See reference \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/ZeroPadding2D\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/ZeroPadding2D)
- Max pooling: [See reference \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/MaxPool2D\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/MaxPool2D)
- Fully connected layer: [See reference \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense)
- Addition: [See reference \(https://www.tensorflow.org/api_docs/python/tf/keras/layers/Add\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Add)

```

In [10]: # GRADED FUNCTION: ResNet50

def ResNet50(input_shape = (64, 64, 3), classes = 6):
    """
    Stage-wise implementation of the architecture of the popular ResNet50:
    CONV2D -> BATCHNORM -> RELU -> MAXPOOL -> CONVBLOCK -> IDBLOCK*2 -> CONVBLOCK -> IDBLOCK*3
    -> CONVBLOCK -> IDBLOCK*5 -> CONVBLOCK -> IDBLOCK*2 -> AVGPPOOL -> FLATTEN -> DENSE

    Arguments:
    input_shape -- shape of the images of the dataset
    classes -- integer, number of classes

    Returns:
    model -- a Model() instance in Keras
    """

    # Define the input as a tensor with shape input_shape
    X_input = Input(input_shape)

    # Zero-Padding
    X = ZeroPadding2D((3, 3))(X_input)

    # Stage 1
    X = Conv2D(64, (7, 7), strides = (2, 2), kernel_initializer = glorot_uniform(seed=0))(X)
    X = BatchNormalization(axis = 3)(X)
    X = Activation('relu')(X)
    X = MaxPooling2D((3, 3), strides=(2, 2))(X)

    # Stage 2
    X = convolutional_block(X, f = 3, filters = [64, 64, 256], s = 1)
    X = identity_block(X, 3, [64, 64, 256])
    X = identity_block(X, 3, [64, 64, 256])

    # YOUR CODE STARTS HERE

    ## Stage 3 (~4 lines)
    X = convolutional_block(X, f = 3, filters = [128,128,512], s = 2)
    X = identity_block(X, 3, [128,128,512])
    X = identity_block(X, 3, [128,128,512])
    X = identity_block(X, 3, [128,128,512])

    ## Stage 4 (~6 lines)
    X = convolutional_block(X, f = 3, filters = [256, 256, 1024], s = 2)
    X = identity_block(X, 3, [256, 256, 1024])
    X = identity_block(X, 3, [256, 256, 1024])
    X = identity_block(X, 3, [256, 256, 1024])
    X = identity_block(X, 3, [256, 256, 1024])
    X = identity_block(X, 3, [256, 256, 1024])

    ## Stage 5 (~3 lines)
    X = convolutional_block(X, f = 3, filters = [512, 512, 2048], s = 2)
    X = identity_block(X, 3, [512, 512, 2048])
    X = identity_block(X, 3, [512, 512, 2048])

    ## AVGPPOOL (~1 line). Use "X = AveragePooling2D(...)(X)"
    X = AveragePooling2D((2, 2))(X)

    # YOUR CODE ENDS HERE

    # output layer
    X = Flatten()(X)
    X = Dense(classes, activation='softmax', kernel_initializer = glorot_uniform(seed=0))(X)

    # Create model
    model = Model(inputs = X_input, outputs = X)

    return model

```

Run the following code to build the model's graph. If your implementation is incorrect, you'll know it by checking your accuracy when running `model.fit(...)` below.

```
In [11]: model = ResNet50(input_shape = (64, 64, 3), classes = 6)
         print(model.summary())
```


Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 64, 64, 3)]	0	
zero_padding2d (ZeroPadding2D)	(None, 70, 70, 3)	0	input_1[0][0]
conv2d_22 (Conv2D)	(None, 32, 32, 64)	9472	zero_padding2d[0][0]
batch_normalization_22 (Batch Normalization)	(None, 32, 32, 64)	256	conv2d_22[0][0]
activation_21 (Activation)	(None, 32, 32, 64)	0	batch_normalization_22[0][0]
max_pooling2d (MaxPooling2D)	(None, 15, 15, 64)	0	activation_21[0][0]
conv2d_23 (Conv2D)	(None, 15, 15, 64)	4160	max_pooling2d[0][0]
batch_normalization_23 (Batch Normalization)	(None, 15, 15, 64)	256	conv2d_23[0][0]
activation_22 (Activation)	(None, 15, 15, 64)	0	batch_normalization_23[0][0]
conv2d_24 (Conv2D)	(None, 15, 15, 64)	36928	activation_22[0][0]
batch_normalization_24 (Batch Normalization)	(None, 15, 15, 64)	256	conv2d_24[0][0]
activation_23 (Activation)	(None, 15, 15, 64)	0	batch_normalization_24[0][0]
conv2d_25 (Conv2D)	(None, 15, 15, 256)	16640	activation_23[0][0]
conv2d_26 (Conv2D)	(None, 15, 15, 256)	16640	max_pooling2d[0][0]
batch_normalization_25 (Batch Normalization)	(None, 15, 15, 256)	1024	conv2d_25[0][0]
batch_normalization_26 (Batch Normalization)	(None, 15, 15, 256)	1024	conv2d_26[0][0]
add_7 (Add)	(None, 15, 15, 256)	0	batch_normalization_25[0][0] batch_normalization_26[0][0]
activation_24 (Activation)	(None, 15, 15, 256)	0	add_7[0][0]
conv2d_27 (Conv2D)	(None, 15, 15, 64)	16448	activation_24[0][0]
batch_normalization_27 (Batch Normalization)	(None, 15, 15, 64)	256	conv2d_27[0][0]
activation_25 (Activation)	(None, 15, 15, 64)	0	batch_normalization_27[0][0]
conv2d_28 (Conv2D)	(None, 15, 15, 64)	36928	activation_25[0][0]
batch_normalization_28 (Batch Normalization)	(None, 15, 15, 64)	256	conv2d_28[0][0]
activation_26 (Activation)	(None, 15, 15, 64)	0	batch_normalization_28[0][0]
conv2d_29 (Conv2D)	(None, 15, 15, 256)	16640	activation_26[0][0]
batch_normalization_29 (Batch Normalization)	(None, 15, 15, 256)	1024	conv2d_29[0][0]
add_8 (Add)	(None, 15, 15, 256)	0	batch_normalization_29[0][0] activation_24[0][0]
activation_27 (Activation)	(None, 15, 15, 256)	0	add_8[0][0]
conv2d_30 (Conv2D)	(None, 15, 15, 64)	16448	activation_27[0][0]
batch_normalization_30 (Batch Normalization)	(None, 15, 15, 64)	256	conv2d_30[0][0]
activation_28 (Activation)	(None, 15, 15, 64)	0	batch_normalization_30[0][0]
conv2d_31 (Conv2D)	(None, 15, 15, 64)	36928	activation_28[0][0]
batch_normalization_31 (Batch Normalization)	(None, 15, 15, 64)	256	conv2d_31[0][0]
activation_29 (Activation)	(None, 15, 15, 64)	0	batch_normalization_31[0][0]
conv2d_32 (Conv2D)	(None, 15, 15, 256)	16640	activation_29[0][0]
batch_normalization_32 (Batch Normalization)	(None, 15, 15, 256)	1024	conv2d_32[0][0]
add_9 (Add)	(None, 15, 15, 256)	0	batch_normalization_32[0][0] activation_27[0][0]
activation_30 (Activation)	(None, 15, 15, 256)	0	add_9[0][0]
conv2d_33 (Conv2D)	(None, 8, 8, 128)	32896	activation_30[0][0]
batch_normalization_33 (Batch Normalization)	(None, 8, 8, 128)	512	conv2d_33[0][0]
activation_31 (Activation)	(None, 8, 8, 128)	0	batch_normalization_33[0][0]
conv2d_34 (Conv2D)	(None, 8, 8, 128)	147584	activation_31[0][0]
batch_normalization_34 (Batch Normalization)	(None, 8, 8, 128)	512	conv2d_34[0][0]
activation_32 (Activation)	(None, 8, 8, 128)	0	batch_normalization_34[0][0]
conv2d_35 (Conv2D)	(None, 8, 8, 512)	66048	activation_32[0][0]
conv2d_36 (Conv2D)	(None, 8, 8, 512)	131584	activation_30[0][0]
batch_normalization_35 (Batch Normalization)	(None, 8, 8, 512)	2048	conv2d_35[0][0]
batch_normalization_36 (Batch Normalization)	(None, 8, 8, 512)	2048	conv2d_36[0][0]
add_10 (Add)	(None, 8, 8, 512)	0	batch_normalization_35[0][0] batch_normalization_36[0][0]
activation_33 (Activation)	(None, 8, 8, 512)	0	add_10[0][0]
conv2d_37 (Conv2D)	(None, 8, 8, 128)	65664	activation_33[0][0]
batch_normalization_37 (Batch Normalization)	(None, 8, 8, 128)	512	conv2d_37[0][0]
activation_34 (Activation)	(None, 8, 8, 128)	0	batch_normalization_37[0][0]
conv2d_38 (Conv2D)	(None, 8, 8, 128)	147584	activation_34[0][0]
batch_normalization_38 (Batch Normalization)	(None, 8, 8, 128)	512	conv2d_38[0][0]
activation_35 (Activation)	(None, 8, 8, 128)	0	batch_normalization_38[0][0]
conv2d_39 (Conv2D)	(None, 8, 8, 512)	66048	activation_35[0][0]

batch_normalization_39 (BatchNo	(None, 8, 8, 512)	2048	conv2d_39[0][0]
add_11 (Add)	(None, 8, 8, 512)	0	batch_normalization_39[0][0] activation_33[0][0]
activation_36 (Activation)	(None, 8, 8, 512)	0	add_11[0][0]
conv2d_40 (Conv2D)	(None, 8, 8, 128)	65664	activation_36[0][0]
batch_normalization_40 (BatchNo	(None, 8, 8, 128)	512	conv2d_40[0][0]
activation_37 (Activation)	(None, 8, 8, 128)	0	batch_normalization_40[0][0]
conv2d_41 (Conv2D)	(None, 8, 8, 128)	147584	activation_37[0][0]
batch_normalization_41 (BatchNo	(None, 8, 8, 128)	512	conv2d_41[0][0]
activation_38 (Activation)	(None, 8, 8, 128)	0	batch_normalization_41[0][0]
conv2d_42 (Conv2D)	(None, 8, 8, 512)	66048	activation_38[0][0]
batch_normalization_42 (BatchNo	(None, 8, 8, 512)	2048	conv2d_42[0][0]
add_12 (Add)	(None, 8, 8, 512)	0	batch_normalization_42[0][0] activation_36[0][0]
activation_39 (Activation)	(None, 8, 8, 512)	0	add_12[0][0]
conv2d_43 (Conv2D)	(None, 8, 8, 128)	65664	activation_39[0][0]
batch_normalization_43 (BatchNo	(None, 8, 8, 128)	512	conv2d_43[0][0]
activation_40 (Activation)	(None, 8, 8, 128)	0	batch_normalization_43[0][0]
conv2d_44 (Conv2D)	(None, 8, 8, 128)	147584	activation_40[0][0]
batch_normalization_44 (BatchNo	(None, 8, 8, 128)	512	conv2d_44[0][0]
activation_41 (Activation)	(None, 8, 8, 128)	0	batch_normalization_44[0][0]
conv2d_45 (Conv2D)	(None, 8, 8, 512)	66048	activation_41[0][0]
batch_normalization_45 (BatchNo	(None, 8, 8, 512)	2048	conv2d_45[0][0]
add_13 (Add)	(None, 8, 8, 512)	0	batch_normalization_45[0][0] activation_39[0][0]
activation_42 (Activation)	(None, 8, 8, 512)	0	add_13[0][0]
conv2d_46 (Conv2D)	(None, 4, 4, 256)	131328	activation_42[0][0]
batch_normalization_46 (BatchNo	(None, 4, 4, 256)	1024	conv2d_46[0][0]
activation_43 (Activation)	(None, 4, 4, 256)	0	batch_normalization_46[0][0]
conv2d_47 (Conv2D)	(None, 4, 4, 256)	590080	activation_43[0][0]
batch_normalization_47 (BatchNo	(None, 4, 4, 256)	1024	conv2d_47[0][0]
activation_44 (Activation)	(None, 4, 4, 256)	0	batch_normalization_47[0][0]
conv2d_48 (Conv2D)	(None, 4, 4, 1024)	263168	activation_44[0][0]
conv2d_49 (Conv2D)	(None, 4, 4, 1024)	525312	activation_42[0][0]
batch_normalization_48 (BatchNo	(None, 4, 4, 1024)	4096	conv2d_48[0][0]
batch_normalization_49 (BatchNo	(None, 4, 4, 1024)	4096	conv2d_49[0][0]
add_14 (Add)	(None, 4, 4, 1024)	0	batch_normalization_48[0][0] batch_normalization_49[0][0]
activation_45 (Activation)	(None, 4, 4, 1024)	0	add_14[0][0]
conv2d_50 (Conv2D)	(None, 4, 4, 256)	262400	activation_45[0][0]
batch_normalization_50 (BatchNo	(None, 4, 4, 256)	1024	conv2d_50[0][0]
activation_46 (Activation)	(None, 4, 4, 256)	0	batch_normalization_50[0][0]
conv2d_51 (Conv2D)	(None, 4, 4, 256)	590080	activation_46[0][0]
batch_normalization_51 (BatchNo	(None, 4, 4, 256)	1024	conv2d_51[0][0]
activation_47 (Activation)	(None, 4, 4, 256)	0	batch_normalization_51[0][0]
conv2d_52 (Conv2D)	(None, 4, 4, 1024)	263168	activation_47[0][0]
batch_normalization_52 (BatchNo	(None, 4, 4, 1024)	4096	conv2d_52[0][0]
add_15 (Add)	(None, 4, 4, 1024)	0	batch_normalization_52[0][0] activation_45[0][0]
activation_48 (Activation)	(None, 4, 4, 1024)	0	add_15[0][0]
conv2d_53 (Conv2D)	(None, 4, 4, 256)	262400	activation_48[0][0]
batch_normalization_53 (BatchNo	(None, 4, 4, 256)	1024	conv2d_53[0][0]
activation_49 (Activation)	(None, 4, 4, 256)	0	batch_normalization_53[0][0]
conv2d_54 (Conv2D)	(None, 4, 4, 256)	590080	activation_49[0][0]
batch_normalization_54 (BatchNo	(None, 4, 4, 256)	1024	conv2d_54[0][0]
activation_50 (Activation)	(None, 4, 4, 256)	0	batch_normalization_54[0][0]
conv2d_55 (Conv2D)	(None, 4, 4, 1024)	263168	activation_50[0][0]
batch_normalization_55 (BatchNo	(None, 4, 4, 1024)	4096	conv2d_55[0][0]
add_16 (Add)	(None, 4, 4, 1024)	0	batch_normalization_55[0][0] activation_48[0][0]
activation_51 (Activation)	(None, 4, 4, 1024)	0	add_16[0][0]
conv2d_56 (Conv2D)	(None, 4, 4, 256)	262400	activation_51[0][0]
batch_normalization_56 (BatchNo	(None, 4, 4, 256)	1024	conv2d_56[0][0]
activation_52 (Activation)	(None, 4, 4, 256)	0	batch_normalization_56[0][0]
conv2d_57 (Conv2D)	(None, 4, 4, 256)	590080	activation_52[0][0]

batch_normalization_57 (BatchNo	(None, 4, 4, 256)	1024	conv2d_57[0][0]
activation_53 (Activation)	(None, 4, 4, 256)	0	batch_normalization_57[0][0]
conv2d_58 (Conv2D)	(None, 4, 4, 1024)	263168	activation_53[0][0]
batch_normalization_58 (BatchNo	(None, 4, 4, 1024)	4096	conv2d_58[0][0]
add_17 (Add)	(None, 4, 4, 1024)	0	batch_normalization_58[0][0] activation_51[0][0]
activation_54 (Activation)	(None, 4, 4, 1024)	0	add_17[0][0]
conv2d_59 (Conv2D)	(None, 4, 4, 256)	262400	activation_54[0][0]
batch_normalization_59 (BatchNo	(None, 4, 4, 256)	1024	conv2d_59[0][0]
activation_55 (Activation)	(None, 4, 4, 256)	0	batch_normalization_59[0][0]
conv2d_60 (Conv2D)	(None, 4, 4, 256)	590080	activation_55[0][0]
batch_normalization_60 (BatchNo	(None, 4, 4, 256)	1024	conv2d_60[0][0]
activation_56 (Activation)	(None, 4, 4, 256)	0	batch_normalization_60[0][0]
conv2d_61 (Conv2D)	(None, 4, 4, 1024)	263168	activation_56[0][0]
batch_normalization_61 (BatchNo	(None, 4, 4, 1024)	4096	conv2d_61[0][0]
add_18 (Add)	(None, 4, 4, 1024)	0	batch_normalization_61[0][0] activation_54[0][0]
activation_57 (Activation)	(None, 4, 4, 1024)	0	add_18[0][0]
conv2d_62 (Conv2D)	(None, 4, 4, 256)	262400	activation_57[0][0]
batch_normalization_62 (BatchNo	(None, 4, 4, 256)	1024	conv2d_62[0][0]
activation_58 (Activation)	(None, 4, 4, 256)	0	batch_normalization_62[0][0]
conv2d_63 (Conv2D)	(None, 4, 4, 256)	590080	activation_58[0][0]
batch_normalization_63 (BatchNo	(None, 4, 4, 256)	1024	conv2d_63[0][0]
activation_59 (Activation)	(None, 4, 4, 256)	0	batch_normalization_63[0][0]
conv2d_64 (Conv2D)	(None, 4, 4, 1024)	263168	activation_59[0][0]
batch_normalization_64 (BatchNo	(None, 4, 4, 1024)	4096	conv2d_64[0][0]
add_19 (Add)	(None, 4, 4, 1024)	0	batch_normalization_64[0][0] activation_57[0][0]
activation_60 (Activation)	(None, 4, 4, 1024)	0	add_19[0][0]
conv2d_65 (Conv2D)	(None, 2, 2, 512)	524800	activation_60[0][0]
batch_normalization_65 (BatchNo	(None, 2, 2, 512)	2048	conv2d_65[0][0]
activation_61 (Activation)	(None, 2, 2, 512)	0	batch_normalization_65[0][0]
conv2d_66 (Conv2D)	(None, 2, 2, 512)	2359808	activation_61[0][0]
batch_normalization_66 (BatchNo	(None, 2, 2, 512)	2048	conv2d_66[0][0]
activation_62 (Activation)	(None, 2, 2, 512)	0	batch_normalization_66[0][0]
conv2d_67 (Conv2D)	(None, 2, 2, 2048)	1050624	activation_62[0][0]
conv2d_68 (Conv2D)	(None, 2, 2, 2048)	2099200	activation_60[0][0]
batch_normalization_67 (BatchNo	(None, 2, 2, 2048)	8192	conv2d_67[0][0]
batch_normalization_68 (BatchNo	(None, 2, 2, 2048)	8192	conv2d_68[0][0]
add_20 (Add)	(None, 2, 2, 2048)	0	batch_normalization_67[0][0] batch_normalization_68[0][0]
activation_63 (Activation)	(None, 2, 2, 2048)	0	add_20[0][0]
conv2d_69 (Conv2D)	(None, 2, 2, 512)	1049088	activation_63[0][0]
batch_normalization_69 (BatchNo	(None, 2, 2, 512)	2048	conv2d_69[0][0]
activation_64 (Activation)	(None, 2, 2, 512)	0	batch_normalization_69[0][0]
conv2d_70 (Conv2D)	(None, 2, 2, 512)	2359808	activation_64[0][0]
batch_normalization_70 (BatchNo	(None, 2, 2, 512)	2048	conv2d_70[0][0]
activation_65 (Activation)	(None, 2, 2, 512)	0	batch_normalization_70[0][0]
conv2d_71 (Conv2D)	(None, 2, 2, 2048)	1050624	activation_65[0][0]
batch_normalization_71 (BatchNo	(None, 2, 2, 2048)	8192	conv2d_71[0][0]
add_21 (Add)	(None, 2, 2, 2048)	0	batch_normalization_71[0][0] activation_63[0][0]
activation_66 (Activation)	(None, 2, 2, 2048)	0	add_21[0][0]
conv2d_72 (Conv2D)	(None, 2, 2, 512)	1049088	activation_66[0][0]
batch_normalization_72 (BatchNo	(None, 2, 2, 512)	2048	conv2d_72[0][0]
activation_67 (Activation)	(None, 2, 2, 512)	0	batch_normalization_72[0][0]
conv2d_73 (Conv2D)	(None, 2, 2, 512)	2359808	activation_67[0][0]
batch_normalization_73 (BatchNo	(None, 2, 2, 512)	2048	conv2d_73[0][0]
activation_68 (Activation)	(None, 2, 2, 512)	0	batch_normalization_73[0][0]
conv2d_74 (Conv2D)	(None, 2, 2, 2048)	1050624	activation_68[0][0]
batch_normalization_74 (BatchNo	(None, 2, 2, 2048)	8192	conv2d_74[0][0]
add_22 (Add)	(None, 2, 2, 2048)	0	batch_normalization_74[0][0] activation_66[0][0]
activation_69 (Activation)	(None, 2, 2, 2048)	0	add_22[0][0]

average_pooling2d (AveragePooli	(None, 1, 1, 2048)	0	activation_69[0][0]
flatten (Flatten)	(None, 2048)	0	average_pooling2d[0][0]
dense (Dense)	(None, 6)	12294	flatten[0][0]
=====			
Total params: 23,600,006			
Trainable params: 23,546,886			
Non-trainable params: 53,120			
None			

```
In [12]: from outputs import ResNet50_summary

model = ResNet50(input_shape = (64, 64, 3), classes = 6)

comparator(summary(model), ResNet50_summary)

All tests passed!
```

As shown in the Keras Tutorial Notebook, prior to training a model, you need to configure the learning process by compiling the model.

```
In [13]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

The model is now ready to be trained. The only thing you need now is a dataset!

Let's load your old friend, the SIGNS dataset.




Figure 6: SIGNS dataset

```
In [14]: X_train_orig, Y_train_orig, X_test_orig, Y_test_orig, classes = load_dataset()

# Normalize image vectors
X_train = X_train_orig / 255.
X_test = X_test_orig / 255.

# Convert training and test labels to one hot matrices
Y_train = convert_to_one_hot(Y_train_orig, 6).T
Y_test = convert_to_one_hot(Y_test_orig, 6).T

print ("number of training examples = " + str(X_train.shape[0]))
print ("number of test examples = " + str(X_test.shape[0]))
print ("X_train shape: " + str(X_train.shape))
print ("Y_train shape: " + str(Y_train.shape))
print ("X_test shape: " + str(X_test.shape))
print ("Y_test shape: " + str(Y_test.shape))

number of training examples = 1080
number of test examples = 120
X_train shape: (1080, 64, 64, 3)
Y_train shape: (1080, 6)
X_test shape: (120, 64, 64, 3)
Y_test shape: (120, 6)
```

Run the following cell to train your model on 2 epochs with a batch size of 32. On a CPU, it should take around 5 minutes per epoch.

```
In [15]: model.fit(X_train, Y_train, epochs = 2, batch_size = 32)

Epoch 1/2
34/34 [=====] - 75s 2s/step - loss: 2.0787 - accuracy: 0.4852
Epoch 2/2
34/34 [=====] - 73s 2s/step - loss: 0.7556 - accuracy: 0.7315

Out[15]: <tensorflow.python.keras.callbacks.History at 0x7f87d0384fd0>
```

Expected Output:

```
Epoch 1/2
34/34 [=====] - 85s 3s/step - loss: 1.9359 - accuracy: 0.5046
Epoch 2/2
34/34 [=====] - 85s 3s/step - loss: 0.6868 - accuracy: 0.8074
```

The exact values could not match, but don't worry about that. The important thing that you must see is that the loss value decreases, and the accuracy increases from the first to second epoch.

Let's see how this model (trained on only two epochs) performs on the test set.

```
In [16]: preds = model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

4/4 [=====] - 1s 299ms/step - loss: 1.8678 - accuracy: 0.6750
Loss = 1.8677754402160645
Test Accuracy = 0.675000011920929
```

Expected Output:

Test Accuracy Around 0.6 +- 0.1

For the purposes of this assignment, you've been asked to train the model for just two epochs. You can see that it performs pretty poorly, but that's ok! The online grader will only run your code for a small number of epochs as well. Please go ahead and submit your assignment as is.

After you have finished this official (graded) part of this assignment, you can also optionally train the ResNet for more iterations, if you want. It tends to get much better performance when trained for ~20 epochs, but this does take more than an hour when training on a CPU.

Using a GPU, this ResNet50 model's weights were trained on the SIGNS dataset. You can load and run the trained model on the test set in the cells below. It may take ≈1min to load the model. Have fun!

```
In [17]: pre_trained_model = tf.keras.models.load_model('resnet50.h5')

In [18]: preds = pre_trained_model.evaluate(X_test, Y_test)
print ("Loss = " + str(preds[0]))
print ("Test Accuracy = " + str(preds[1]))

4/4 [=====] - 1s 225ms/step - loss: 0.1596 - accuracy: 0.9500
Loss = 0.15958662331104279
Test Accuracy = 0.949999988079071
```

Congratulations on finishing this assignment! You've now implemented a state-of-the-art image classification system! Woo hoo!

ResNet50 is a powerful model for image classification when it's trained for an adequate number of iterations. Hopefully, from this point, you can use what you've learned and apply it to your own classification problem to perform state-of-the-art accuracy.

What you should remember:

- Very deep "plain" networks don't work in practice because vanishing gradients make them hard to train.
- Skip connections help address the Vanishing Gradient problem. They also make it easy for a ResNet block to learn an identity function.
- There are two main types of blocks: The **identity block** and the **convolutional block**.
- Very deep Residual Networks are built by stacking these blocks together.

5 - Test on Your Own Image (Optional/Ungraded)

If you wish, you can also take a picture of your own hand and see the output of the model. To do this:

1. Click on "File" in the upper bar of this notebook, then click "Open" to go on your Coursera Hub.
2. Add your image to this Jupyter Notebook's directory, in the "images" folder
3. Write your image's name in the following code
4. Run the code and check if the algorithm is right!

```
In [ ]: img_path = 'images/my_image.jpg'
img = image.load_img(img_path, target_size=(64, 64))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = x/255.0
print('Input image shape:', x.shape)
imshow(img)
prediction = pre_trained_model.predict(x)
print("Class prediction vector [p(0), p(1), p(2), p(3), p(4), p(5)] = ", prediction)
print("Class:", np.argmax(prediction))
```

You can also print a summary of your model by running the following code.

```
In [ ]: pre_trained_model.summary()
```

6 - Bibliography

This notebook presents the ResNet algorithm from He et al. (2015). The implementation here also took significant inspiration and follows the structure given in the GitHub repository of Francois Chollet:

- Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun - [Deep Residual Learning for Image Recognition \(2015\) \(https://arxiv.org/abs/1512.03385\)](https://arxiv.org/abs/1512.03385)
- Francois Chollet's GitHub repository: <https://github.com/fchollet/deep-learning-models/blob/master/resnet50.py> (https://github.com/fchollet/deep-learning-models/blob/master/resnet50.py)

```
In [ ]:
```